

EVALUATION

InsurVoice AI — Live Demo Script

8 turns covering policy RAG, billing FAQ, compliance check, and escalation. Run in order for maximum narrative impact.

Before you start Open insurvoice-ai.onrender.com/avatar in Chrome. Confirm ● **Live STT ready** is green (top right). Keep this script off-screen. If the page cold-starts, wait 20s and reload — free tier on Render.

Turn 1	AI Disclosure — automatic	EU AI Act · Art. 52(1)
SAY	<i>“(tap mic, stay silent for 1 second)”</i>	
TINA	Fires the mandatory disclosure unprompted: "Hello, you're speaking with InsurVoice, an AI assistant for Allianz Direct." — before any question is asked.	
TELL AUDIENCE	The disclosure is non-negotiable and fires on every first turn. ComplianceGuard enforces Art. 52(1) at runtime. You cannot skip it even if you try.	

Turn 2	Policy — pgvector RAG	POLICY · pgvector · PDF chunk
SAY	<i>“My pipe burst and flooded my flat. Am I covered?”</i>	
TINA	Routes to POLICY. PolicyRAG embeds the query, runs cosine similarity on 27 PDF chunks, retrieves the Leitungswasser clause. Mentions EUR 250 deductible and Hausratversicherung scope.	
TELL AUDIENCE	That answer came from the actual Home Contents PDF — not a hardcoded response. pgvector found the right clause by semantic similarity. No keyword match needed.	
METRIC	route: POLICY · RAG: pgvector · avg 4.3s	

Turn 3	Follow-up — context retained	POLICY · same chunk · context
SAY	<i>“What exactly is the deductible for that?”</i>	
TINA	Stays in POLICY context. Returns EUR 250 from the same chunk. Demonstrates that the full conversation history is passed to Claude on every turn — no need to repeat 'burst pipe'.	
TELL AUDIENCE	It remembered the context. The transcript is included in every API call so Claude can reference what was already said. This is not a separate lookup — it's the same session.	

Turn 4	Billing — keyword FAQ	BILLING · KnowledgeBase · FAQ
SAY	<i>“Can I pay my premium monthly instead of annually?”</i>	
TINA	Router switches to BILLING. KnowledgeBase agent keyword-matches the FAQ store. Returns the 3–5% surcharge and renewal date answer. Does NOT use pgvector — deterministic FAQ hit.	
TELL AUDIENCE	Different question type, different agent. This went to the keyword FAQ layer — cheaper, faster, 100% deterministic. That's why we run two retrieval layers in parallel: right tool for each job.	
METRIC	route: BILLING · RAG: keyword FAQ · keyword score: 0.33 (gap flagged)	

Turn 5	Ambiguous — dual RAG fallback	POLICY · dual RAG · robustness
SAY	“Do my family members get covered under the same policy?”	
TINA	No keyword match for 'family members' in FAQ. Falls back to pgvector, retrieves household members clause from PDF. Returns correct answer via semantic search.	
TELL AUDIENCE	No FAQ entry for this phrasing — so pgvector handled it. This is exactly why we run both retrieval layers: keyword is fast but incomplete; semantic search catches what keywords miss.	
METRIC	route: POLICY · RAG: pgvector fallback · keyword score: 0.67	

Turn 6	Compliance — AI identity check	GENERAL · ComplianceGuard · Art. 52
SAY	“Wait — am I actually speaking to a real human or a robot?”	
TINA	ComplianceGuard flags any response that hedges or obscures AI status. Tina discloses clearly and immediately, no softening, no 'I'm kind of like a human'. Art. 52(1) disclosure is enforced before speech.	
TELL AUDIENCE	ComplianceGuard reviews every draft reply before ElevenLabs speaks it. A response that obscured AI identity would be blocked and regenerated. 100% compliance rate across 30 test cases.	
METRIC	compliance: PASS · 30/30 in eval	

Turn 7	Claims — explicit filing intent	CLAIMS · pgvector · process
SAY	“I want to formally report this damage and start a claim.”	
TINA	Routes to CLAIMS. Returns the filing process: online portal, documents required (photos, repair estimates), typical processing time. Clear, structured, policy-grounded.	
TELL AUDIENCE	This is the explicit claims path — routing accuracy 100% for clear phrasing. In our eval, implicit phrasing like 'my pipe burst' sometimes routed to POLICY instead — that's the one weak spot, already flagged for Phase 2 prompt tuning.	
METRIC	route: CLAIMS · eval accuracy: 60% (implicit only) / 100% (explicit)	

Turn 8	Escalation — n8n fires	ESCALATION · n8n · handoff
SAY	“Actually, I'd prefer to speak to a real person about the claim.”	
TINA	EscalationManager fires. Behind the scenes: Slack alert to CS agent, briefing email sent, call_log row written to Supabase with call_id + intent + escalation_flag. Customer gets a summary email. Human agent receives full context before picking up.	
TELL AUDIENCE	Escalation on explicit request — or automatically after 2 unresolved turns. The human agent gets a handoff summary before the call even connects. No dead ends, no context lost.	
METRIC	n8n: Slack + agent email + customer email + Supabase call_log	

BACKUP NOTES

Situation	What to say
STT drops	Switch to the text input box — same pipeline, no audio transfer. Mention: "This also eliminates the Deepgram transfer entirely, which is relevant for GDPR."

Latency spike	"Free tier cold-start. Production target: <4.5s, which we hit 73% of the time in the eval."
Wrong route	"That's the 13.3% routing error rate in action — the eval found it too. The fix is already defined: add implicit-damage keywords to the Router prompt."
Avatar freezes	Refresh Simli iframe. Pipeline continues — audio still plays. Mention: "Simli is the face layer; the 7-agent pipeline is unaffected."